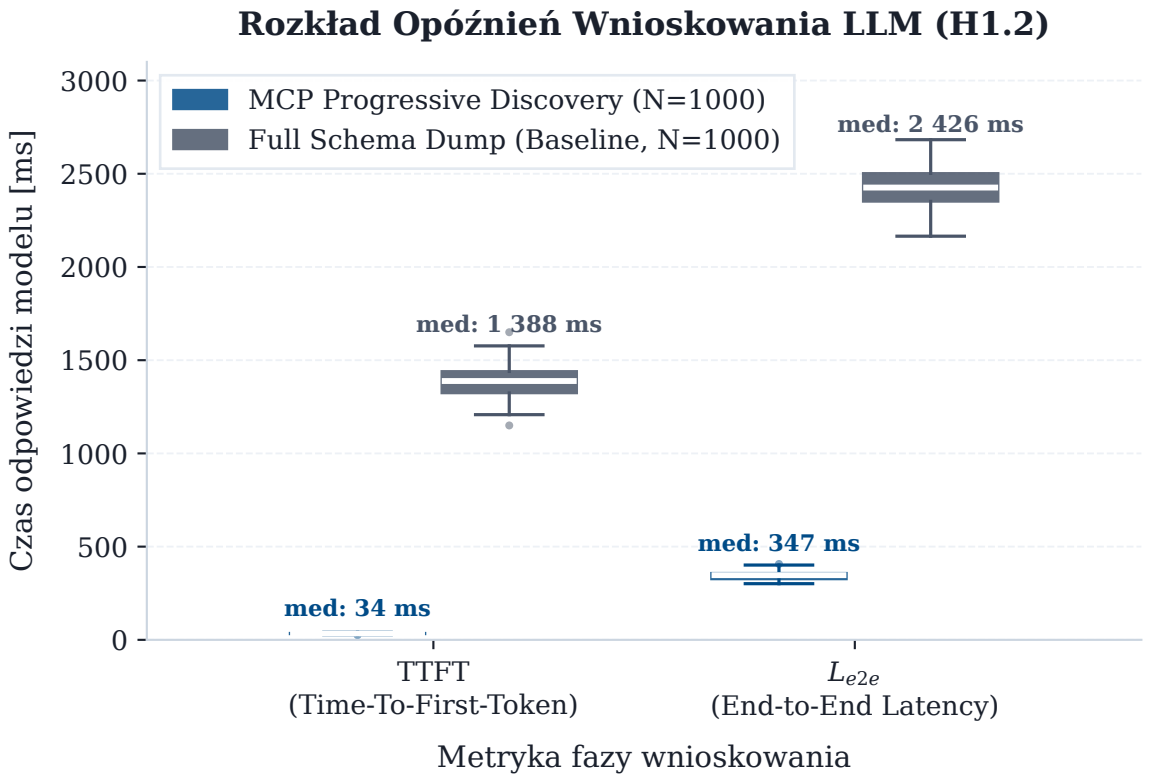
Rozkład Opóźnień Wnioskowania LLM (H1.2)
MCP Progressive Discovery (N=1000)
Full Schema Dump (Baseline, N=1000)
med: 2 426 ms
med: 1 388 ms
med: 347 ms
med: 34 ms
3000
2500
2000
1500
1000
500
0
Czas odpowiedzi modelu [ms]
TTFT
(Time-To-First-Token)
L_{e2e}
(End-to-End Latency)
Metryka fazy wnioskowania